

All are Worth Words: A ViT Backbone for Diffusion Models

Department of Computer Engineering and Artificial Intelligence
MAI Lab.

YS Park



CVPR 2023 open access

These CVPR 2023 papers are the Open Access versions, provided by the [Computer Vision Foundation](#).

Except for the watermark, they are identical to the accepted versions; the final published version of the proceedings is available on IEEE Xplore.

This material is presented to ensure timely dissemination of scholarly and technical work. Copyright and all rights therein are retained by authors or by other copyright holders. All persons copying this information are expected to adhere to the terms and constraints invoked by each author's copyright.

All Are Worth Words: A ViT Backbone for Diffusion Models

Fan Bao, Shen Nie, Kaiwen Xue, Yue Cao, Chongxuan Li, Hang Su, Jun Zhu; Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023, pp. 22669-22679

Abstract


Vision transformers (ViT) have shown promise in various vision tasks while the U-Net based on a convolutional neural network (CNN) remains dominant in diffusion models. We design a simple and general ViT-based architecture (named U-ViT) for image generation with diffusion models. U-ViT is characterized by treating all inputs including the time, condition and noisy image patches as tokens and employing long skip connections between shallow and deep layers. We evaluate U-ViT in unconditional and class-conditional image generation, as well as text-to-image generation tasks, where U-ViT is comparable if not superior to a CNN-based U-Net of a similar size. In particular, latent diffusion models with U-ViT achieve record-breaking FID scores of 2.29 in class-conditional image generation on ImageNet 256x256, and 5.48 in text-to-image generation on MS-COCO, among methods without accessing large external datasets during the training of generative models. Our results suggest that, for diffusion-based image modeling, the long skip connection is crucial while the down-sampling and up-sampling operators in CNN-based U-Net are not always necessary. We believe that U-ViT can provide insights for future research on backbones in diffusion models and benefit generative modeling on large scale cross-modality datasets.



- Introduction
- Method
- Related Work
- Experiments
- Conclusion

- Diffusion models emerge recently for high quality image generation.
 - Applications: text-to-image, image-to-image, video generations, etc
 - And they employ U-Net(CNN) as a backbone.
- On the other hand ViT have shown promise in various vision task.
 - Therefore, a very natural question arises: whether the reliance of the CNN-based U-Net is necessary in diffusion models?
- This paper introduces simple ViT-based architecture called **U-ViT**.

- U-ViT is a simple and general backbone for diffusion models in image generation.
 - They also takes time t , condition c , noised image x_t (as tokens).
 - The long skip connections provide shortcuts for the low-level features.
- U-ViT optionally adds a 3×3 convolutional block before output.
 - This is intended to prevent the potential artifacts in images produced by transformers



- **The way to combine the long skip branch**
 - **(1) Linear(Concat(h_m h_s))** (2) $h_m + h_s$ (3) $h_m + \text{Linear}(h_s)$ (4) $\text{Linear}(h_m + h_s)$ (5) dropped
- **The way to feed the time into the network**
 - **(1) treat is as a token** (2) incorporate the time after the layer norm in the transformer block.
- **The way to add an extra convolutional block after the transformer**
 - **(1) add a 3×3 conv. after the linear proj.** (2) add a 3×3 conv. before linear proj (3) dropped
- **Variants of the patch embedding**
 - **(1) The original patch embedding adopts a linear proj. that maps a patch to a token embedding.**
 - (2) use a stack of 3×3 conv. followed by a 1×1 conv. to map an image to token embeddings.
- **Variants of the position embedding**
 - **(1) 1D** (2) 2D







Figure 2. Ablate design choices. The one marked with * is the adopted choice of U-ViT illustrated in Figure 1. Since this ablation aims to determine implementation details, we evaluate FID on 10K generated samples (instead of 50K samples for efficiency).



(a) Depth (#layers)



(b) Width (hidden size)



(c) Patch size

Figure 3. Effect of depth, width and patch size. The one marked with * corresponds to the setting of U-ViT-S/2 (see Table 2).

- Transformers in diffusion models
 - GenViT: long skip connection / incorporates time before norm. layers
 - VQ-Diffusion: obtains a sequence of discrete image tokens via a VQ-GAN
- U-Net in diffusion models
 - initially introduce CNN-based U-Net to model the gradient of log-likelihood function for continuous image data.
 - After that, improvements on the CNN-based U-Net: group norm., multi-head attention, improved residual block and cross attention.
- Improvements of diffusion models
 - Fast sampling, improved training, controllable generation.

- Datasets.
 - For unconditional learning
 - CIFAR10 which contain 50K training images
 - CelebA 64×64, which contain 162,770 training images of human faces.
 - For class-conditional learning
 - ImageNet 64×64, 256×256 and 512×512 resolutions, which contains 1,281,167 training images from 1K different classes.
 - For text-to-image learning
 - MS-COCO at 256×256 resolution, which contains 82,783 training images and 40,504 validation images. Each image is annotated with 5 captions.
- High resolution image generation.
 - LDM for images at 256×256 and 512×512 resolutions. We firstly convert them to latent representations at 32×32 and 64×64 resolutions respectively, using a pretrained image autoencoder provided by Stable Diffusion.
 - Then we model these latent representations using the proposed U-ViT.
- Text-to-image learning.
 - On MS-COCO, we convert discrete texts to a sequence of embeddings using a CLIP text encoder following Stable Diffusion. Then these embeddings are fed into U-ViT as a sequence of tokens.

- High resolution image generation.
 - LDM for images at 256×256 and 512×512 resolutions. We firstly convert them to latent representations at 32×32 and 64×64 resolutions respectively, using a pretrained image autoencoder provided by Stable Diffusion.
 - Then we model these latent representations using the proposed U-ViT.
- Text-to-image learning.
 - On MS-COCO, we convert discrete texts to a sequence of embeddings using a CLIP text encoder following Stable Diffusion. Then these embeddings are fed into U-ViT as a sequence of tokens.

Model	#Layers	Hidden size D	MLP size	#Heads	#Params
U-ViT-Small	13	512	2048	8	44M
U-ViT-Small (Deep)	17	512	2048	8	58M
U-ViT-Mid	17	768	3072	12	131M
U-ViT-Large	21	1024	4096	16	287M
U-ViT-Huge	29	1152	4608	16	501M

Table 2. Configurations of U-ViT.

- Training.
 - Optimizer:
 - AdamW with a weight decay of 0.3 for all datasets.
 - learning rate:
 - $2e-4$: most datasets
 - $3e-4$: ImageNet 64×64
 - Iterations:
 - 500K
 - CIFAR10 and CelebA 64×64 with a batch size of 128.
 - ImageNet 512×512 , with a batch size of 1024.
 - 300K
 - ImageNet 64×64 and ImageNet 256×256 .
 - 1M
 - MS-COCO with a batch size of 256.
 - classifier-free guidance:
 - ImageNet 256×256 , ImageNet 512×512 and MS-COCO

Model on CIFAR10		FID ↓
GAN		
StyleGAN2-ADA [31]		2.92
Diff. based on U-Net	#Params	
DDPM [25]	36M	3.17
DDPM [49]	53M	2.90
DDPM++ cont. [67]	62M	2.55
EDM [†] [30]	56M	1.97
Diff. based on ViT	#Params	
GenViT [76]	11M	20.20
U-ViT-S/2	44M	3.11
Model on CelebA 64×64		
FID ↓		
Diff. based on U-Net	#Params	
DDIM [63]	79M	3.26
Soft Truncation [†] [32]	62M	1.90
Diff. model based on ViT	#Params	
U-ViT-S/4	44M	2.87
Model on ImageNet 64×64		
FID ↓		
GAN		
BigGAN-deep [5]		4.06
StyleGAN-XL [59]		1.51
Diff. based on U-Net	#Params	
DDPM (small) [49]	100M	6.92
DDPM (large) [49]	270M	2.92
CDM [26]	Unknown	1.48
ADM [13]	296M	2.07
EDM [†] [30]	296M	1.36
Diff. based on ViT	#Params	
U-ViT-M/4	131M	5.85
U-ViT-L/4	287M	4.26

Model on ImageNet 256×256		FID ↓
GAN		
BigGAN-deep [5]		6.95
StyleGAN-XL [59]		2.30
Discrete diff. based on transformer		
VQ-Diffusion [20]		11.89
VQ-Diffusion (acc0.05) [20]		5.32
Diff. based on U-Net	#Params	
DDPM [49]	270M + 280M (SR)	12.26
CDM [26]	Unknown	4.88
ADM [13]	554M	10.94
ADM-U [13]	296M + 312M (SR)	7.49
ADM-G [13]	554M + 54M (Cls)	4.59
ADM-G, ADM-U [13]	296M + 65M (Cls) + 312M (SR)	3.94
LDM [‡] [54]	400M + 55M (AE)	3.60
Diff. based on ViT	#Params	
U-ViT-H/2 [‡]	501M + 84M (AE)	2.29
Model on ImageNet 512×512		
FID ↓		
GAN		
BigGAN-deep [5]		8.43
StyleGAN-XL [59]		2.41
Diff. based on U-Net	#Params	
ADM [13]	559M	23.24
ADM-U [13]	422M + 309M (SR)	9.96
ADM-G [13]	559M + 54M (Cls)	7.72
ADM-G, ADM-U [13]	422M + 43M (Cls) + 309M (SR)	3.85
Diff. based on ViT	#Params	
U-ViT-H/4 [‡]	501M + 84M (AE)	4.05

	4	5	10	15	20
LDM (trained 178K)	34.48	12.73	4.51	3.87	3.68
U-ViT-H/2 (trained 200K)	16.48	4.94	3.87	3.54	2.91
U-ViT-H/2 (trained 500K)	15.44	4.64	3.18	2.92	2.53

Table 3. FID results on ImageNet 256×256 under different number of sampling steps using the DPM-Solver sampler [44].




Figure 5. Ablate the long skip connection on ImageNet 256×256 (w/o classifier-free guidance).

Table 1. FID results of unconditional image generation on CIFAR10 and CelebA 64×64, and class-conditional image generation on ImageNet 64×64, 256×256 and 512×512. We mark the best results *among diffusion models* in **bold**. We also include GAN results (gray) for completeness. Methods marked with [†] use advanced training techniques for diffusion models. Methods marked with [‡] model latent representations of images [54] and use classifier-free guidance [27]. We also present the number of parameters of auxiliary components for diffusion models, where SR represents a super-resolution module, AE represents an image autoencoder, and Cls represents a classifier.




Figure 4. Image generation results of U-ViT: selected samples on ImageNet 512×512 and ImageNet 256×256 , and random samples on CIFAR10, CelebA 64×64 , and ImageNet 64×64 .

Model	FID	Type	Training datasets	#Params
Generative model trained on external large dataset (zero-shot)				
DALL-E [53]	~ 28	Autoregressive	DALL-E dataset (250M)	12B
CogView [14]	27.1	Autoregressive	Internal dataset (30M)	4B
LAFITE [82]	26.94	GAN	CC3M (3M)	75M + 151M (TE)
GLIDE [48]	12.24	Diffusion	DALL-E dataset (250M)	3.5B + 1.5B (SR)
Make-A-Scene [19]	11.84	Autoregressive	Union datasets (without MS-COCO) (35M)	4B
DALL-E 2 [52]	10.39	Diffusion	DALL-E dataset (250M)	4.5B + 700M (SR)
Imagen [56]	7.27	Diffusion	Internal dataset (460M) + LAION (400M)	2B + 4.6B (TE) + 600M (SR)
Parti [77]	7.23	Autoregressive	LAION (400M) + FIT (400M) + JFT (4B)	20B + 630M (AE)
Re-Imagen [8]	6.88	Diffusion	KNN-ImageText (50M)	2.5B + 750M (SR)
Generative model trained on external large dataset with access to MS-COCO				
VQ-Diffusion [†] [20]	13.86	Discrete diffusion	Conceptual Caption Subset (7M)	370M
Make-A-Scene [19]	7.55	Autoregressive	Union datasets (with MS-COCO) (35M)	4B
Re-Imagen [‡] [8]	5.25	Diffusion	KNN-ImageText (50M)	2.5B + 750M (SR)
Parti [†] [77]	3.22	Autoregressive	LAION (400M) + FIT (400M) + JFT (4B)	20B + 630M (AE)
Generative model trained on MS-COCO				
AttnGAN [75]	35.49	GAN	MS-COCO (83K)	230M
DM-GAN [83]	32.64	GAN	MS-COCO (83K)	46M
VQ-Diffusion [20]	19.75	Discrete diffusion	MS-COCO (83K)	370M
DF-GAN [70]	19.32	GAN	MS-COCO (83K)	19M
XMC-GAN [79]	9.33	GAN	MS-COCO (83K)	166M
Friro [18]	8.97	Diffusion	MS-COCO (83K)	512M + 186M (TE) + 68M (AE)
LAFITE [82]	8.12	GAN	MS-COCO (83K)	75M + 151M (TE)
U-Net*	7.32	Latent diffusion	MS-COCO (83K)	53M + 123M (TE) + 84M (AE)
U-ViT-S/2	5.95	Latent diffusion	MS-COCO (83K)	45M + 123M (TE) + 84M (AE)
U-ViT-S/2 (Deep)	5.48	Latent diffusion	MS-COCO (83K)	58M + 123M (TE) + 84M (AE)

Table 4. FID results of different models on MS-COCO validation (256×256). U-ViT-S (Deep) increases the number of layers from 13 to 17 compared to U-ViT-S. We also present the number of parameters of auxiliary components for a model when it is reported in the corresponding paper, where SR represents a super-resolution module, AE represents an image autoencoder, and TE represents a text encoder. Methods marked with [†] finetune on MS-COCO. Methods marked with [‡] use MS-COCO as a knowledge base for retrieval. The U-Net* is trained by ourselves to serve as a direct baseline of U-ViT, where we leave other parts unchanged except for the backbone.

Experiments/Text-to-Image Generation on MS-COCO



Figure 6. Text-to-image generation on MS-COCO. All the other settings except the backbone are the same. U-Net and U-ViT generate samples using the same random seed for a fair comparison. The random seed is unselected.

- This work presents U-ViT, a simple and general ViT-based architecture for image generation with diffusion models.
- U-ViT treats all inputs including the time, condition and noisy image patches as tokens and employs long skip connections between shallow and deep layers.
- We evaluate U-ViT in various tasks.
- U-ViT is comparable if not superior to a CNN-based U-Net of a similar size.
- These results suggest that, for diffusion-based image modeling, the long skip connection is crucial while the down/up-sampling operators in CNN-based U-Net are not always necessary.